

## Handwritten Digit Recognizer

```

1 #Step 1: Install & Import Libraries
2 import numpy as np
3 import tensorflow as tf
4 from tensorflow import keras
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
7 from tensorflow.keras.utils import to_categorical
8 import matplotlib.pyplot as plt
9 import seaborn as sns

```

```

1 #Step 2: Load and Preprocess the Dataset
2 # Load MNIST dataset
3 from tensorflow.keras.datasets import mnist
4 (X_train, y_train), (X_test, y_test) = mnist.load_data()
5 # Reshape data to match CNN input format (28x28x1)
6 X_train = X_train.reshape(-1, 28, 28, 1).astype('float32') / 255.0
7 X_test = X_test.reshape(-1, 28, 28, 1).astype('float32') / 255.0
8 # Convert labels to one-hot encoding
9 y_train = to_categorical(y_train, 10)
10 y_test = to_categorical(y_test, 10)

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
 11490434/11490434 ————— 0s 0us/step

```

1 #Step 3: Define the CNN Model
2 # Define CNN architecture
3 model = Sequential([
4 Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(28,28,1)),
5 MaxPooling2D(pool_size=(2,2)),
6 Conv2D(64, kernel_size=(3,3), activation='relu'),
7 MaxPooling2D(pool_size=(2,2)),
8 Flatten(),
9 Dense(128, activation='relu'),
10 Dropout(0.5),
11 Dense(10, activation='softmax')
12 ])
13 # Compile the model
14 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape` to `input\_shape` in the constructor of `Conv2D` or `Conv3D` layers. This is deprecated and will be removed in a future version. Please use the `input\_shape` argument in the `compile` method instead.  
 super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```

1 #Step 4: Train the Model
2 # Train the model
3 history = model.fit(X_train, y_train, epochs=10, batch_size=128, validation_data=(X_test, y_test))

```

Epoch 1/10  
 469/469 ————— 46s 95ms/step - accuracy: 0.8095 - loss: 0.6031 - val\_accuracy: 0.9789 - val\_loss: 0.0635  
 Epoch 2/10  
 469/469 ————— 80s 91ms/step - accuracy: 0.9698 - loss: 0.1037 - val\_accuracy: 0.9866 - val\_loss: 0.0401  
 Epoch 3/10  
 469/469 ————— 82s 90ms/step - accuracy: 0.9785 - loss: 0.0725 - val\_accuracy: 0.9883 - val\_loss: 0.0344  
 Epoch 4/10  
 469/469 ————— 81s 88ms/step - accuracy: 0.9822 - loss: 0.0591 - val\_accuracy: 0.9898 - val\_loss: 0.0307  
 Epoch 5/10  
 469/469 ————— 83s 90ms/step - accuracy: 0.9840 - loss: 0.0509 - val\_accuracy: 0.9898 - val\_loss: 0.0283  
 Epoch 6/10  
 469/469 ————— 81s 89ms/step - accuracy: 0.9863 - loss: 0.0429 - val\_accuracy: 0.9904 - val\_loss: 0.0284  
 Epoch 7/10  
 469/469 ————— 41s 88ms/step - accuracy: 0.9897 - loss: 0.0349 - val\_accuracy: 0.9910 - val\_loss: 0.0273  
 Epoch 8/10  
 469/469 ————— 42s 90ms/step - accuracy: 0.9888 - loss: 0.0353 - val\_accuracy: 0.9927 - val\_loss: 0.0218  
 Epoch 9/10  
 469/469 ————— 82s 90ms/step - accuracy: 0.9909 - loss: 0.0296 - val\_accuracy: 0.9921 - val\_loss: 0.0236  
 Epoch 10/10  
 469/469 ————— 83s 91ms/step - accuracy: 0.9918 - loss: 0.0256 - val\_accuracy: 0.9908 - val\_loss: 0.0253

```

1 #Step 5: Evaluate Model Performance
2 # Evaluate on test data
3 loss, accuracy = model.evaluate(X_test, y_test)

```

```

4 print(f"Test Loss: {loss}")
5 print(f"Test Accuracy: {accuracy * 100:.2f}%")


```

313/313 ————— 3s 9ms/step - accuracy: 0.9888 - loss: 0.0323  
 Test Loss: 0.025310376659035683  
 Test Accuracy: 99.08%

```

1 #Step 6: Visualize Training Results
2 # Plot accuracy trends
3 plt.plot(history.history['accuracy'], label='Train Accuracy')
4 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
5 plt.xlabel('Epochs')
6 plt.ylabel('Accuracy')
7 plt.legend()
8 plt.show()

```



```

1 #Step 7: Make Predictions & Visualize Results
2 import numpy as np
3 def predict_digit(index):
4     img = X_test[index].reshape(1, 28, 28, 1)
5     prediction = model.predict(img)
6     predicted_label = np.argmax(prediction)
7     actual_label = np.argmax(y_test[index])
8     plt.imshow(X_test[index].reshape(28, 28), cmap='gray')
9     plt.title(f"Predicted: {predicted_label}, Actual: {actual_label}")
10    plt.show()
11 # Test prediction
12 predict_digit(0)
13

```

↻ 1/1 — 0s 107ms/step

Predicted: 7, Actual: 7

